Statistical Machine Learning Project 2

Arnav Chakravarthy
ASU ID: 1216873944

Required Tasks

The specific algorithmic tasks I needed to perform for this part of the project included:

- Initialize clusters using Strategy 1 and 2. Strategy 1 involves randomly picking the initial centers from the given samples. Strategy 2 involves picking only the first point randomly and picking points with ith center (i>1) such that the average distance of this chosen point to all previous centroids is maximal.

- Test the implementation with the number k of clusters ranging from 2-10.

- Plot the objective function value vs. the number of clusters k. Under each strategy, plot the objective function twice, each start from a different initialization.

1. Initialization

1.1. Strategy 1

I have used the "NumPy" library for the Python programming language. It adds support for large, multi-dimensional arrays and matrices.

And for this whole project I have used vectorization to make calculations more efficient.

Each sample is represented as a 2D Vector.

I have used the random.sample to choose "k" random and unique points to avoid duplicate clusters.

I have then stored this in a Python list called all_centers. The shape of this would be (k x 2) Where k is the number of centroids.

1.2. Strategy 2

For this strategy, I choose the first point randomly using random.randint.

For the rest of the centroids, I built a function called *calc_point_distances()* which calculates the average distance of each point to the previously calculated centroids, stores them in a list in decreasing order and returns them.

I now iterate through this list until I find the first index which is not already present in *all_clusters*. Hence I calculate the point with the maximum average distance which is not already present in *all_clusters*. The reason for doing this and not just returning an argmax is because after considering a certain number of centers, the future centers get repeated. Hence this following method gives us unique centroids.

2. K-Means Algorithm

I created a Python dictionary called *center_points_dict* which stores the cluster index as the key, and the points present in the cluster as its value. The key is an integer and the value is a list of lists.

For every iteration of KMeans these points get updated as well as all_clusters get updated.

I have used Euclidean distance to assign a point to a cluster such that the Euclidean distance between the point and the centroid of the assigned cluster is the least among that of all the other centroids.

After each iteration I update all_clusters by updating each centroid as the average of all the points present in that specific cluster. These points can be retrieved from *center_points_dict*.

I also created two plotting functions to visualize the centroids as well as all the points in their respective clusters after performing K-Means.

K=7, Strategy 1

K=7, Strategy=2

3. Objective Function

$$\sum_{i=1}^{k} \sum_{x \epsilon D_i} ||x - \mu_i||^2$$

k – number of clusters

D – Total Sample set

D_i - Samples in cluster i

μ_i – Centroid of cluster i.

4. Plotting the Objective Function vs Number of Graphs

I used matplotlib as my plotting library of choice.

I first created a function called *compute_all_objective_costs()*, which calculated the final objective function cost after convergence for clusters ranging from 2-10, and also take the strategy as an input. Since I created KMeans as a class, I just had to pass in number of clusters, strategy type, maximum number of iterations and the dataset path as parameters to the class instance.

I was then able to get the final objective cost for each instance each time with different number of clusters and a specific strategy time.

These costs were stored in a list and returned.

I then used matplotlib to plot these costs on the Y-axis and the number of clusters on the X-axis.

Strategy 1

Random Seed 80

Random Seed 20

Strategy 2

Random Seed 20

Random Seed 80